



ELSI Interface (Development Version) User's Guide

The ELSI Team
elsi-interchange.org

April 9, 2018

Contents

1	Introduction	3
1.1	The Cubic Wall of Kohn-Sham Density-Functional Theory	3
1.2	ELSI, the ELectronic Structure Infrastructure	4
1.2.1	Design Tenets of ELSI	4
1.3	Kohn-Sham Solver Libraries Supported by ELSI	5
1.3.1	ELPA: Eigenvalue soLvers for Petaflop-Applications	5
1.3.2	libOMM: Orbital Minimization Method	5
1.3.3	PEXSI: Pole EXpansion and Selected Inversion	5
1.4	Acknowledgments	6
2	Installation of ELSI	7
2.1	Overview	7
2.2	Prerequisites	7
2.3	CMake Basics	7
2.4	Configuration	8
2.4.1	Compilers	8
2.4.2	Solvers	9
2.4.3	Build Targets	9
2.4.4	“Toolchain” Files	10
2.5	Importing ELSI into KS-DFT Codes	10
2.5.1	Linking a KS-DFT Code against ELSI	10
2.5.2	Importing ELSI into a KS-DFT Code	10
3	The ELSI API	11
3.1	Overview of the ELSI API	11
3.2	Description of All ELSI APIs	11
3.2.1	Initializing ELSI	11
3.2.2	Setting Up MPI	13
3.2.3	Setting Up Matrix Formats	13
3.2.4	Setting Up Multiple k-points and/or Spin Channels	14
3.2.5	Finalizing ELSI	15
3.3	Solving Eigenvalues and Eigenvectors	15
3.4	Computing Density Matrices	16
3.5	Customizing ELSI	17
3.5.1	Customizing the ELSI Interface	17
3.5.2	Customizing the ELPA Solver	19
3.5.3	Customizing the libOMM Solver	19
3.5.4	Customizing the PEXSI Solver	19
3.6	Getting Additional Results from ELSI	21
3.6.1	Getting Results from the ELSI Interface	22
3.6.2	Getting Results from the PEXSI Solver	22
3.7	Demonstration Pseudo-Code	23
3.7.1	2D Block-Cyclic Distributed Dense Matrix + ELSI Eigensolver Interface	23
3.7.2	1D Block Distributed CSC Sparse Matrix + ELSI Eigensolver Interface	23
3.7.3	1D Block-Cyclic Distributed CSC Sparse Matrix + ELSI Eigensolver Interface	24
3.7.4	2D Block-Cyclic Distributed Dense Matrix + ELSI Density Matrix Interface	25

3.7.5	1D Block Distributed CSC Sparse Matrix + ELSI Density Matrix Interface	25
3.7.6	1D Block-Cyclic Distributed CSC Sparse Matrix + ELSI Density Matrix Interface	26
3.7.7	Multiple k-points Calculations	27
3.8	C/C++ Interface	27

1 Introduction

1.1 The Cubic Wall of Kohn-Sham Density-Functional Theory

Molecular and materials simulations based on Kohn-Sham density-functional theory (KS-DFT) [1] are widely used to provide atomic-scale insights, understanding, and predictions across a wide range of disciplines in the sciences and in engineering.

In KS-DFT [1], the many-electron problem for the Born-Oppenheimer electronic ground state is reduced to a system of single particle equations known as the Kohn-Sham equations

$$\hat{h}^{\text{KS}}\psi_l = \epsilon_l\psi_l, \quad (1.1)$$

where ψ_l and ϵ_l are Kohn-Sham orbitals and their associated eigenenergies, and \hat{h}^{KS} denotes the Kohn-Sham Hamiltonian. In practice, N_{basis} basis functions $\phi_i(\mathbf{r})$ are employed to expand the Kohn-Sham orbitals:

$$\psi_l(\mathbf{r}) = \sum_{j=1}^{N_{\text{basis}}} c_{jl}\phi_j(\mathbf{r}). \quad (1.2)$$

Using non-orthogonal basis functions (e.g., Gaussian functions, Slater functions, numeric atom-centered orbitals) in Eq. 1.2 converts Eq. 1.1 to a generalized eigenvalue problem that can be expressed in the following matrix form:

$$\mathbf{H}\mathbf{c} = \mathbf{\epsilon}\mathbf{S}\mathbf{c}. \quad (1.3)$$

Here, \mathbf{H} and \mathbf{S} are the Hamiltonian matrix and the overlap matrix, whose elements can be computed through numerical integrations. The matrix \mathbf{c} and diagonal matrix $\mathbf{\epsilon}$ contain the eigenvectors and eigenvalues, respectively, of the eigensystem.

When using orthonormal basis sets (e.g., plane waves, multi-resolution wavelets), the eigenproblem described in Eq. 1.3 reduces to a standard form where \mathbf{S} is an identity matrix.

The explicit solution of Eq. 1.3 yields the Kohn-Sham orbitals ψ_l , from which the electron density $n(\mathbf{r})$ can be computed following an orbital-based method that scales as $O(N^2)$:

$$n(\mathbf{r}) = \sum_{j=1}^{N_{\text{basis}}} f_l \psi_l^*(\mathbf{r}) \psi_l(\mathbf{r}), \quad (1.4)$$

where f_l denotes the occupation number of each orbital. In an actual computation, it is sufficient to perform the summation only for the occupied ($f_l > 0$) orbitals. The ratio of occupied orbitals to the total number of basis functions can be below 1% for plane wave basis sets, whereas with some localized basis sets, fewer basis functions are required, leading to a larger fraction of occupied states typically between 10% and 40%.

An alternative method that scales as $O(N)$ can be employed for localized basis functions:

$$n(\mathbf{r}) = \sum_{i,j}^{N_{\text{basis}}} \phi_i^*(\mathbf{r}) n_{ij} \phi_j(\mathbf{r}), \quad (1.5)$$

with n_{ij} being the elements of the density matrix that need to be computed before the density update.

From a viewpoint of computational complexity, almost all standard pieces of solving the Kohn-Sham equations can be formulated in a linear scaling fashion with respect to the system size. The only piece that can not be easily addressed in an $O(N)$ fashion is the Kohn-Sham eigenvalue problem.

Traditional solutions to the Kohn-Sham eigenvalue problem rely on diagonalization of the Hamiltonian matrix to obtain the Kohn-Sham orbitals ψ_l . The associated computational cost is $O(N^3)$, regardless of the diagonalization algorithm employed. Thanks to a low prefactor, these methods are advantageous to use for systems comprised of up to roughly a few thousands of atoms, which account for the bulk of KS-DFT applications. Systems beyond a few thousands of atoms, however, cannot be efficiently handled by diagonalization due to its cubic scaling nature. Alternatively, approaches with a lower scaling factor, e.g. linear scaling methods, usually directly compute the density matrix as a function of the Hamiltonian matrix. The large prefactor of these methods introduces a significant overhead for small systems, hindering the establishment of lower scaling approaches as mainstream methods of the field. In addition, such alternatives are typically restricted to certain classes of systems or problems. The transition from cubic scaling methods to linear scaling ones is therefore not trivial to automate. In practice, the KS eigenvalue problem remains to be a bottleneck of KS-DFT simulations on current HPC architectures and for system sizes significantly exceeding several thousands of atoms.

1.2 ELSI, the ELectronic Structure Infrastructure

ELSI unifies the community effort in overcoming the cubic-wall problem of KS-DFT by bridging the divide between developers of electronic structure solvers and KS-DFT codes. Via a unified interface, ELSI gives KS-DFT developers easy access to multiple solvers that solve or circumvent the Kohn-Sham eigenproblem efficiently. Solvers are treated on equal footing within ELSI, giving solver developers a unified platform for implementation and benchmarking across codes and physical systems. Solvers may be switched dynamically in an SCF cycle, allowing the KS-DFT developer to mix-and-match strengths of different solvers. Solvers can work cooperatively with one another within ELSI, allowing for acceleration greater than either solver can achieve individually. Most importantly, ELSI exists as a community for KS-DFT and solver developers to interact and work together to improve performance of solvers, with monthly web meetings to discuss progress on code development, yearly on-site “connector meetings”, and planned webinars and workshops.

The current development version of ELSI supports ELPA [3, 4], libOMM [2], and PEXSI [5, 6] solvers. Codes currently integrated with ELSI include DFTB+ [7], DGDFT [8], FHI-aims [9], NWChem [10] via Global Arrays, and SIESTA [11].

1.2.1 Design Tenets of ELSI

Versatility: ELSI supports real-valued and complex-valued density matrix, eigenvalue, and eigenvector calculations. A unified software interface designed for rapid integration into a variety of KS-DFT codes is provided. Fortran and C/C++ interfaces are provided.

Flexibility: ELSI supports both dense and sparse matrices as input/output. Supported matrix distribution layouts include 2D block-cyclic distribution, 1D block-cyclic distribution, and 1D block distribution. In situations where the input/output matrix format requested by the calling KS-DFT code and the format used internally by the requested solver are different, conversion and redistribution of matrices will be performed automatically.

Scalability: The solver libraries collected in ELSI are highly scalable. For instance, ELPA can scale to $\sim 100k$ CPU cores given a sufficiently large problem to solve, and PEXSI, with its efficient two-level parallelism, easily scales to $\sim 500k$ CPU cores.

Portability: ELSI and its redistributed library source packages have been confirmed to work on commonly-used HPC architectures (Cray, IBM, Intel, NVIDIA) using major compilers (Cray, GNU, IBM, Intel, PGI).

1.3 Kohn-Sham Solver Libraries Supported by ELSI

The current development version of ELSI stably supports ELPA [3, 4], libOMM [2], and PEXSI [5, 6] solvers. Since this user’s guide focuses on the ELSI interface layer, given below are very brief introductions of the solvers. For more details and technical descriptions of each solver, the reader is referred to the original publications, e.g., those cited here, as well as an overview paper of the 2017.05 version of ELSI [12].

1.3.1 ELPA: Eigenvalue solvers for Petaflop-Applications

ELPA [3, 4] solves the Kohn-Sham eigenvalue problem in Eq. 1.3 by direct (tri)diagonalization.

In ELPA, the generalized eigenproblem is first transformed to the standard form by using Cholesky factorization of the overlap matrix. Then, the standard eigenproblem is either directly reduced to the tridiagonal form, using similar algorithms as implemented in traditional ScaLAPACK eigensolvers, or first reduced to a banded intermediate form, then to the tridiagonal form. After the eigenvectors of the resulted tridiagonal system are solved, they must be back-transformed to the original form in either one step or two steps, depending on the algorithm employed in the tridiagonalization. Compared to the traditional one-step approach, although the two-step alternative introduces two additional steps (one forward transformation and one backward), it has been shown to enable faster computation and better parallel scalability on present-day computers. Specifically, the matrix-vector operations (BLAS level-2 routines) in the one-step tridiagonalization can be mostly replaced by more efficient matrix-matrix operations (BLAS level-3 routines) in the two-step version of the algorithm. The additional effort needed by back-transforming the eigenvectors is often alleviated in electronic structure calculations due to the fact that only a relatively small fraction of the eigensolutions need to be retrieved.

Since ELPA employs the same 2D block-cyclic matrix distribution as does ScaLAPACK, it can easily be substituted into existing codes that already support parallel linear algebra by ScaLAPACK.

1.3.2 libOMM: Orbital Minimization Method

Instead of direct diagonalization, the orbital minimization method (OMM) relies on efficient iterative algorithms to minimize an unconstrained energy functional using a set of auxiliary Wannier functions that expand the occupied subspace of the system. The OMM energy functional, when minimized, can be shown to be equal to the sum of the eigenvalues of the occupied Kohn-Sham orbitals. Furthermore, the Wannier functions are driven towards perfect orthonormality at this minimum. The density matrix is then constructed from the final Wannier functions. Although this density matrix is sufficient for the electron density update following Eq. 1.5, without knowledge of individual eigenstates, the OMM approach cannot handle systems with fractional occupation numbers resulting, e.g., from a finite electronic temperature, such as is typically required for metals.

Compared to other minimization methods with the orthonormality constraint of eigenstates, the advantage of OMM is that it only requires an unconstrained minimization without an explicit orthonormalization step. This makes the OMM a good candidate for linear scaling DFT. However, in order to do so, it is necessary to spatially confine the Wannier functions, which introduces a number of technical difficulties. Recently, the OMM algorithm has been found to result in an efficient iterative solver with conventional cubic scaling but a smaller prefactor than diagonalization. This has been taken by the new implementation in libOMM [2]. The minimization of the OMM energy functional is carried out in libOMM by using the conjugate-gradient (CG) method, with an efficient preconditioning using the kinetic energy matrix when possible.

1.3.3 PEXSI: Pole EXpansion and Selected Inversion

The pole expansion and selected inversion (PEXSI) method [5, 6] provides another alternative way for solving the Kohn-Sham electronic structure without diagonalization. As a Fermi operator expansion (FOE) based method, PEXSI expands the density matrix using a P -term pole expansion, whose number of terms of the pole expansion is proportional to $\log(\beta\Delta E)$, where $\beta = 1/(k_B T)$ is the inverse of the thermal energy and ΔE is the spectral radius. This logarithmic scaling makes the pole expansion a highly efficient approach to expand the Fermi operator. Typically only ~ 20 poles are sufficient for the result obtained from PEXSI to be fully comparable to that obtained from diagonalization.

In PEXSI, selected elements of the density matrix, corresponding to the non-zero pattern of \mathbf{H} and \mathbf{S} , are evaluated using a parallel selected inversion algorithm. The computational cost of the PEXSI technique depends on the dimensionality of the system: $O(N)$ for 1D systems, $O(N^{1.5})$ for 2D systems, and $O(N^2)$ for general 3D bulk systems. This favorable scaling hinges on the sparse character of the Hamiltonian and overlap matrices, but not on any fundamental assumption about the localization properties of the single particle density matrix, i.e., this method is applicable to metallic systems as well as insulating and semi-conducting ones.

The PEXSI method has a multi-level parallelism structure and is by design highly scalable. The recently developed massively parallel PEXSI technique can make efficient use of $10,000 \sim 100,000$ processors on high performance machines.

1.4 Acknowledgments

ELSI is a National Science Foundation Software Infrastructure for Sustained Innovation - Scientific Software Integration (SI2-SSI) supported software infrastructure project. The ELSI Interface software and this User's Guide are based upon work supported by the National Science Foundation under Grant Number 1450280. Any opinions, findings, and conclusions or recommendations expressed here are those of the authors and do not necessarily reflect the views of the National Science Foundation.

2 Installation of ELSI

2.1 Overview

The ELSI package contains the ELSI Interface software as well as redistributed source code for the solver libraries ELPA (version 2017.05.001), libOMM, and PEXSI (version 1.0.0). We highly encourage all users to request access to our GitLab server at elsi-team@duke.edu, where we regularly update ELSI between releases while preserving stability.

Starting from February 2018, the installation of ELSI makes use of the CMake software.

2.2 Prerequisites

To build ELSI, the minimum requirements are:

CMake	[version 3.0 or newer]
Fortran compiler	[with Fortran 2003]
C compiler	[with C99]
MPI	

The following two additional components enable the PEXSI solver, thus are highly recommended:

C++ compiler	[with C++ 11]
PtScotch library	[version 6.0.4 or newer]

To build ELSI as a shared library, or to build the ELSI test programs, linear algebra libraries should be provided for ELSI to link against:

[BLAS](#), [LAPACK](#), [BLACS](#), [ScaLAPACK](#)

By default, solver libraries redistributed within ELSI will be built. Optionally, they may be substituted by user's optimized versions:

ELPA	[version 2017.05.003 or newer]
libOMM	
MatrixSwitch	
SuperLU_DIST	[version 5.1.3 or newer]

2.3 CMake Basics

The typical workflow of using CMake to build ELSI looks like:


```

$ ls

CMakeLists.txt  external/  src/  test/  ...

$ mkdir build
$ cd build
$ cmake .. -DCMAKE_INSTALL_PREFIX=/path/to/install/elsi [-Doptions ...]

...
...
-- Generating done
-- Build files have been written to: /current/dir

$ make [-j np]
$ make install

```

whenever CMake is invoked, one of the command line arguments must point to the path where the top level CMakeLists.txt file exists, hence the “..” in the above example.

An option can be defined by adding

```
-DKeyword=Value
```

to the command line when invoking CMake. If “Keyword” is of type boolean, its “Value” may be “ON” or “OFF”. If “Keyword” defines a list of libraries or include directories, the items should be separated with “;” (semicolon).

For example,

```

-DCMAKE_C_COMPILER=gcc
-DENABLE_TESTS=OFF
-DENABLE_PEXSI=ON
-DMATH_LIB="/path/to/libscalapack.a;/path/to/lapack.a;/path/to/blas.a"
-DELPA_INC="/path/to/include;/another/path/to/include"

```

Available options are introduced in the next section.

2.4 Configuration

2.4.1 Compilers

CMake automatically detects and sets default compilers. The choices made by CMake often work, but not necessarily guarantee the optimal performance. In some cases, the compilers picked up by CMake may not be the ones desired by the user. Therefore, it is mandatory that the user sets the identification of the compiler explicitly:

```

-DCMAKE_Fortran_COMPILER=YOUR_MPI_FORTRAN_COMPILER
-DCMAKE_C_COMPILER=YOUR_MPI_C_COMPILER
-DCMAKE_CXX_COMPILER=YOUR_MPI_C++_COMPILER

```

In addition, it is highly recommended to specify the compiler flags, in particular the optimization flags:

```

-DCMAKE_Fortran_FLAGS=YOUR_FORTRAN_COMPILE_FLAGS
-DCMAKE_C_FLAGS=YOUR_MPI_C_COMPILE_FLAGS
-DCMAKE_CXX_FLAGS=YOUR_MPI_C++_COMPILE_FLAGS

```

Please note that the C++ compiler is not needed if building ELSI without PEXSI.

2.4.2 Solvers

ELPA, libOMM, PEXSI, and SuperLU_DIST are redistributed with the current ELSI package. The internal version of ELPA comes with a few “kernels” specifically written to take advantage of processor architecture (e.g. vectorization instruction set extensions). A kernel may be chosen by the “ELPA2_KERNEL” keyword. Available options are:

```
-DELPA2_KERNEL=BGQ
-DELPA2_KERNEL=AVX
-DELPA2_KERNEL=AVX2
-DELPA2_KERNEL=AVX512
```

In ELPA, these kernels are employed to accelerate the calculation of eigenvectors, which is often a computational bottleneck when a large percentage of eigenvectors need to be calculated. If this is the case in the user’s application, it is highly recommended that the user consults their system administrators and selects the kernel most suited to their system architecture.

Experienced users are encouraged to link the ELSI interface against external, better optimized solver libraries. Relevant options for this purpose are:

```
-DELPA_LIB=YOUR_ELPA_LIBRARY(IES)
-DELPA_INC=YOUR_ELPA_INCLUDE_PATH(S)
-DOmm_LIB=YOUR_LIBOMM_AND_MATRIXSWITCH_LIBRARIES
-DOmm_INC=YOUR_LIBOMM_AND_MATRIXSWITCH_INCLUDE_PATH(S)
-DENABLE_PEXSI=ON
-DSUPERLU_LIB=YOUR_SUPERLU_DIST_LIBRARY(IES)
-DSUPERLU_INC=YOUR_SUPERLU_DIST_INCLUDE_PATH(S)
-DPTSCOTCH_LIB=YOUR_PTSCOTCH_LIBRARY(IES)
-DPTSCOTCH_INC=YOUR_PTSCOTCH_INCLUDE_PATH(S)
```

If both “XXX_LIB” and “XXX_INC” are valid, the internal version of “XXX” library will not be built.

Please note that in the current version of ELSI, PEXSI is not enabled by default, and an external version of PEXSI is not officially supported. When enabling PEXSI, the user should have a working PtScotch library, as PtScotch is not redistributed through ELSI.

2.4.3 Build Targets

By default, a static library (libelsi.a) will be created as the target of the compilation. Since no shared library or executable is generated, there is no need to provide linear algebra libraries, i.e., BLAS, LAPACK, BLACS, and ScaLAPACK. Building ELSI as a shared library may be enabled by:

```
-DBUILD_SHARED_LIBS=ON
```

Building ELSI test programs may be enabled by:

```
-DENABLE_TESTS=ON
```

In either case, the linear algebra libraries mentioned above must be valid through

```
-DMATH_LIB=YOUR_LINEAR_ALGEBRA_LIBRARY(IES)
```

If the test programs are enabled, the installation may be verified by

```
$ ...
$ make [-j np]
$ make test
```

or

```
$ ...
$ make [-j np]
$ ctest
```

Note that the tests may not run if launching MPI jobs is prohibited on the user's working platform.

2.4.4 “Toolchain” Files

It is sometimes convenient to edit the settings in a “toolchain” file that can be read by CMake:

```
-DCMAKE_TOOLCHAIN_FILE=YOUR_TOOLCHAIN_FILE
```

Example “toolchains” are provided in the “./toolchains” directory of the ELSI package, which the user may use as templates to create new ones.

2.5 Importing ELSI into KS-DFT Codes

2.5.1 Linking a KS-DFT Code against ELSI

ELSI places its library and include files in the `lib` and `include` subfolders, respectively, of the ELSI root directory. An example set of compiler flags to link a generic KS-DFT code against a standard installation of ELSI are:

```
ELSI_INCLUDE = -I/PATH/TO/BUILD/ELSI/include
ELSI_LIB      = -L/PATH/TO/BUILD/ELSI/lib -lelsi \
               -l0MM -lMatrixSwitch -lelpa -lpexsi -lsuperlu_dist \
               -L$/PATH/TO/SCOTCH/lib -lptscotchparmetis -lptscotch \
               -lptscotcherr -lscotchmetis -lscotch -lscotcherr
```

These flags may vary based on the user's needs. Linking ELSI against preinstalled solver libraries will require the user modify these flags accordingly.

If PEXSI is not enabled, an example set of the resulting compiler flags would be:

```
ELSI_INCLUDE = -I/PATH/TO/BUILD/ELSI/include
ELSI_LIB      = -L/PATH/TO/BUILD/ELSI/lib -lelsi \
               -l0MM -lMatrixSwitch -lelpa
```

2.5.2 Importing ELSI into a KS-DFT Code

ELSI may be used in the KS-DFT code by importing the appropriate header file. For codes written in Fortran, this is done by using the ELSI module

```
USE ELSI
```

For codes written in C, the ELSI wrapper may be imported by including the header file

```
#include <elsi.h>
```

These import statements give the KS-DFT code access to the ELSI interface. In the next chapter, we will describe the API for the ELSI interface.

3 The ELSI API

3.1 Overview of the ELSI API

In this chapter, we present the public-facing API for the ELSI Interface. We anticipate that fine details of this interface may change slightly in the future, but the fundamental structure of the interface layer is expected to remain consistent.

ELSI provides a C interface in addition to the native Fortran interface. The vast majority of this chapter, while written from a Fortran-ic standpoint, applies equally to both interfaces. Information specifically about the C wrapper for ELSI may be found in Section 3.8.

To allow multiple instances of ELSI to co-exist within a single calling code, we define an `elsi_handle` data type to encapsulate the state of an ELSI instance, i.e., all runtime parameters associated with the ELSI instance.

An `elsi_handle` instance is initialized with the `elsi_init` subroutine and is subsequently passed to all other ELSI subroutine calls.

3.2 Description of All ELSI APIs

In this section, we introduce the interface subroutines of ELSI. While this section serves as a reference to all subroutines, the user is encouraged to explore the demonstration codes of the ELSI interface in Section 3.7.

3.2.1 Initializing ELSI

The ELSI interface must be initialized via the `elsi_init` subroutine before any other ELSI subroutine may be called.

`elsi_init`(handle, solver, parallel_mode, matrix_format, n_basis, n_electron, n_state)

Argument	Data Type	in/out	Explanation
<code>handle</code>	<code>type(elsi_handle)</code>	out	Handle to the current ELSI instance.
<code>solver</code>	integer	in	The desired KS solver. Accepted values are: 1 (ELPA), 2 (LIBOMM), and 3 (PEXSI).
<code>parallel_mode</code>	integer	in	The type of parallelization used. See comment 3. Accepted values are: 0 (SINGLE_PROC) and 1 (MULTI_PROC).
<code>matrix_format</code>	integer	in	The matrix format used for the matrices handled by ELSI (Hamiltonian, overlap, density matrix, etc.) See comment 1. Accepted values are: 0 (BLACS_DENSE), 1 (PEXSI_CSC), and 2 (SIESTA_CSC).
<code>n_basis</code>	integer	in	Number of basis functions, i.e. the global size of the Hamiltonian matrix.
<code>n_electron</code>	real double	in	Number of electrons.
<code>n_state</code>	integer	in	Number of states, as interpreted by the solver. See comment 2.

Comments

1) **matrix_format**: **BLACS_DENSE**(0) refers a dense matrix format in a 2-dimensional block-cyclic distribution, i.e. the BLACS standard. **PEXSI_CSC**(1) refers to a compressed sparse column (CSC) matrix format in a 1-dimensional block distribution. **SIESTA_CSC**(2) refers to a compressed sparse column (CSC) matrix format in a 1-dimensional block-cyclic distribution. As the Hamiltonian, overlap, and density matrices are symmetric (Hermitian), compressed sparse row (CSR) matrix format is effectively supported.

2) **n.state**: If ELPA is the chosen solver, this parameter specifies the number of eigenstates to solve by the eigensolver. If libOMM is the chosen solver, **n.state** must be exactly the number of occupied states, as libOMM cannot handle fractional occupation numbers[2]. PEXSI does not make use of this parameter and a dummy value may be passed.

3) **parallel_mode**: The two allowed values of **parallel_mode**, 0 (**SINGLE_PROC**) and 1 (**MULTI_PROC**), allow for three parallelization strategies commonly employed by KS-DFT codes. See below.

3a) **SINGLE_PROC**: Solves the KS eigenproblem following a LAPACK-like fashion. This option may only be selected when ELPA is chosen as the solver. This allows the following parallelization strategy:

SINGLE_PROC Example:

Every MPI task independently handles a group of k-points uniquely assigned to it.

Example number of k-points: 16

Example number of MPI tasks: 4

MPI task 0 handles k-points 1, 2, 3, 4 sequentially;

MPI task 1 handles k-points 5, 6, 7, 8 sequentially;

MPI task 2 handles k-points 9, 10, 11, 12 sequentially;

MPI task 3 handles k-points 13, 14, 15, 16 sequentially.

Pseudocode 1:

```
elsi_init(elsi_h, ..., parallel_mode=0, ...)
...
for i_kpt = 1, n_kpoints_local, 1 do
    elsi_ev_{real|complex}(elsi_h, ham_this_kpt, ovlp_this_kpt, eval_this_kpt, evec_this_kpt)
end
```

3b) **MULTI_PROC**: Solves the KS eigenproblem following a ScaLAPACK-like fashion. This allows the usage of the following two parallelization strategies:

MULTI_PROC Example:

Groups of MPI tasks coordinate to handle the same k-point, uniquely assigned to that group.

Example number of k-points: 4

Example number of MPI tasks: 16

MPI tasks 0, 1, 2, 3 cooperatively handle k-point 1;

MPI tasks 4, 5, 6, 7 cooperatively handle k-point 2;

MPI tasks 8, 9, 10, 11 cooperatively handle k-point 3;

MPI tasks 12, 13, 14, 15 cooperatively handle k-point 4.

Pseudocode 2:

```
elsi_init(elsi_h, ..., parallel_mode=1, ...)
elsi_set_mpi(elsi_h, my_mpi_comm)
...
elsi_ev_{real|complex}(elsi_h, my_ham, my_ovlp, my_eval, my_evec)

or

elsi_init(elsi_h, ..., parallel_mode=1, ...)
elsi_set_mpi(elsi_h, my_mpi_comm)
elsi_set_kpoint(elsi_h, n_kpt, my_kpt, my_weight)
elsi_set_mpi_global(elsi_h, mpi_comm_global)
...
elsi_dm_complex(elsi_h, my_ham, my_ovlp, my_dm, global_energy)
```

In the last example, please note that when there is more than one k-point, a global MPI communicator must be provided for inter-k-points communications. See the documentation of [elsi_set_kpoint](#), [elsi_set_spin](#), and [elsi_set_mpi_global](#).

3.2.2 Setting Up MPI

The MPI communicator used by ELSI is passed into ELSI by the calling code via the [elsi_set_mpi](#) subroutine. When there is more than one k-point and/or spin channel, this communicator will be used only for one task defined

```
elsi_set_mpi(handle, mpi_comm)
```

Argument	Data Type	in/out	Explanation
handle	type(elsi_handle)	inout	Handle to the current ELSI instance.
mpi_comm	integer	in	MPI communicator.

3.2.3 Setting Up Matrix Formats

When using the BLACS matrix format (BLACS_DENSE), BLACS parameters are passed into ELSI via the [elsi_set_blacs](#) subroutine. The matrix format used internally in the ELSI interface and the ELPA solver requires the block sizes of the 2-dimensional block-cyclic distribution are the same in the row and column directions. It is necessary to call this subroutine before calling any solver interface that makes use of BLACS_DENSE matrix format.

```
elsi_set_blacs(handle, blacs_context, block_size)
```

Argument	Data Type	in/out	Explanation
handle	type(elsi_handle)	inout	Handle to the current ELSI instance.
blacs_context	integer	in	BLACS context.
block_size	integer	in	Block size of the 2D block-cyclic distribution, specifying both row and column directions.

When using the sparse matrix formats, namely 1D block distributed compressed sparse column (PEXSI_CSC) or 1D block-cyclic distributed compressed sparse column (SIESTA_CSC), the sparsity pattern should be passed into ELSI via the [elsi_set_csc](#) subroutine. It is necessary to call this subroutine before calling any solver interface that makes use of the sparse matrix formats.

`elsi_set_csc(handle, global_nnz, local_nnz, local_col, row_index, col_pointer)`

Argument	Data Type	in/out	Explanation
<code>handle</code>	<code>type(elsi_handle)</code>	inout	Handle to the current ELSI instance.
<code>global_nnz</code>	integer	in	Global number of non-zeros.
<code>local_nnz</code>	integer	in	Local number of non-zeros.
<code>local_col</code>	integer	in	Local number of matrix columns.
<code>col_pointer</code>	integer, rank-1 array	in	Local column pointer array. Dimension: <code>local_nnz</code> .
<code>row_index</code>	integer, rank-1 array	in	Local row index array. Dimension: <code>local_col+1</code> .

When using the 1D block distributed compressed sparse column (PEXSI_CSC) format, the block size of the 1D distribution cannot be set by the user. This is because the PEXSI solver requires that the block size must be `n_basis / n_procs`, where `n_basis` is the total number of basis functions, and `n_procs` is the number of MPI tasks assigned to this problem. When using the 1D block-cyclic distributed compressed sparse column (SIESTA_CSC) format, the block size of the 1D distribution must be set by the user.

`elsi_set_csc_blk(handle, block_size)`

Argument	Data Type	in/out	Explanation
<code>handle</code>	<code>type(elsi_handle)</code>	inout	Handle to the current ELSI instance.
<code>global_nnz</code>	integer	in	Block size of the 1D block-cyclic distribution.

3.2.4 Setting Up Multiple k-points and/or Spin Channels

When there is more than one k-point and/or spin channel in the simulating system, the ELSI interface can be configured to support parallel treatment of the k-points and/or spin channels. This is achieved by calling `elsi_set_kpoint` and `elsi_set_spin` to pass the required information into ELSI, and calling `elsi_set_mpi_global` to provide ELSI a global MPI communicator, which is used for a few communications among all the k-points and spin channels. Note that the current ELSI interface only supports the case where the eigenproblems for all the k-points and spin channels are fully parallelized, i.e., there is no MPI task that handles more than one k-point and/or spin channel.

`elsi_set_kpoint(handle, n_kpt, i_kpt, weight)`

Argument	Data Type	in/out	Explanation
<code>handle</code>	<code>type(elsi_handle)</code>	inout	Handle to the current ELSI instance.
<code>n_kpt</code>	integer	in	Total number of k-points.
<code>i_kpt</code>	integer	in	Index of the k-point handled by this MPI task.
<code>weight</code>	integer	in	Weight of the k-point handled by this MPI task.

`elsi_set_spin(handle, n_spin, i_spin)`

Argument	Data Type	in/out	Explanation
<code>handle</code>	<code>type(elsi_handle)</code>	inout	Handle to the current ELSI instance.
<code>n_spin</code>	integer	in	Total number of spin channels.
<code>i_spin</code>	integer	in	Index of the spin channel handled by this MPI task.

`elsi_set_mpi_global(handle, mpi_comm_global)`

Argument	Data Type	in/out	Explanation
<code>handle</code>	<code>type(elsi_handle)</code>	inout	Handle to the current ELSI instance.
<code>mpi_comm_global</code>	integer	in	Global MPI communicator used for communications among all k-points and spin channels.

3.2.5 Finalizing ELSI

When an ELSI instance is no longer needed, its associated handle should be cleaned up by calling `elsi_finalize`.

`elsi_finalize(handle)`

Argument	Data Type	in/out	Explanation
<code>handle</code>	<code>type(elsi_handle)</code>	inout	Handle to the current ELSI instance.

3.3 Solving Eigenvalues and Eigenvectors

The following subroutines return all the eigenvalues and (a subset of) eigenvectors of the provided H and S matrices. Only eigensolvers may be selected as the desired solver when using these subroutines.

`elsi_ev_real(handle, ham, ovlp, eval, evec)`

Argument	Data Type	in/out	Explanation
<code>handle</code>	<code>type(elsi_handle)</code>	inout	Handle to the current ELSI instance.
<code>ham</code>	real double, rank-2 array	inout	Real Hamiltonian matrix in 2D block-cyclic dense format. See comment 1.
<code>ovlp</code>	real double, rank-2 array	inout	Real overlap matrix (or its Cholesky factor) in 2D block-cyclic dense format. See comment 1.
<code>eval</code>	real double, rank-1 array	inout	Eigenvalues. See comment 2.
<code>evec</code>	real double, rank-2 array	out	Real eigenvectors in 2D block-cyclic dense format. See comment 3.

`elsi_ev_complex(handle, ham, ovlp, eval, evec)`

Argument	Data Type	in/out	Explanation
<code>handle</code>	<code>type(elsi_handle)</code>	inout	Handle to the current ELSI instance.
<code>ham</code>	complex double, rank-2 array	inout	Complex Hamiltonian matrix in 2D block-cyclic dense format. See comment 1.
<code>ovlp</code>	complex double, rank-2 array	inout	Complex overlap matrix (or its Cholesky factor) in 2D block-cyclic dense format. See comment 1.
<code>eval</code>	real double, rank-1 array	inout	Eigenvalues. See comments 2.
<code>evec</code>	complex double, rank-2 array	out	Complex eigenvectors in 2D block-cyclic dense format. See comment 3.

`elsi_ev_real_sparse(handle, ham, ovlp, eval, evec)`

Argument	Data Type	in/out	Explanation
<code>handle</code>	<code>type(elsi_handle)</code>	inout	Handle to the current ELSI instance.
<code>ham</code>	real double, rank-1 array	inout	Real Hamiltonian matrix in 1D block or block-cyclic CSC sparse format.
<code>ovlp</code>	real double, rank-1 array	inout	Real overlap matrix in 1D block or block-cyclic CSC sparse format.
<code>eval</code>	real double, rank-1 array	inout	Eigenvalues. See comment 2.
<code>evec</code>	real double, rank-2 array	out	Real eigenvectors in 2D block-cyclic dense format. See comment 3.

`elsi_ev_complex_sparse(handle, ham, ovlp, eval, evec)`

Argument	Data Type	in/out	Explanation
<code>handle</code>	<code>type(elsi_handle)</code>	inout	Handle to the current ELSI instance.
<code>ham</code>	complex double, rank-1 array	inout	Complex Hamiltonian matrix in 1D block or block-cyclic CSC sparse format.
<code>ovlp</code>	complex double, rank-1 array	inout	Complex overlap matrix in 1D block or block-cyclic CSC sparse format.
<code>eval</code>	real double, rank-1 array	inout	Eigenvalues. See comment 2.
<code>evec</code>	complex double, rank-2 array	out	Complex eigenvectors in 2D block-cyclic dense format. See comment 3.

Comments

- 1) The Hamiltonian matrix will be destroyed by ELPA during computation. ELPA will overwrite the overlap matrix in its initial execution with the Cholesky factor, which will be reused by subsequent subroutine calls to `elsi_ev_real` or `elsi_ev_complex`. When using `elsi_ev_real_sparse`, the Cholesky factor (which is not sparse) is stored internally in the BLACS_DENSE format. Starting from the second call to `elsi_ev_real_sparse`, the input sparse overlap matrix will not be referenced.
- 2) When using the ELPA solver, `elsi_ev_real`, `elsi_ev_complex`, `elsi_ev_real_sparse`, and `elsi_ev_complex_sparse` always compute all the eigenvalues, regardless of the choice of `n_state` specified in `elsi_init`. The dimension of `eval` thus should always be `n_basis`.
- 3) When using the ELPA solver, `elsi_ev_real`, `elsi_ev_complex`, `elsi_ev_real_sparse`, and `elsi_ev_complex_sparse` compute a subset of all eigenvectors. The number of eigenvectors to compute is specified by the keyword `n_state` in `elsi_init`. However, the local `eigenvectors` array should always be initialized to correspond to a global array of size `n_basis × n_basis`, whose extra part is used as working space in ELPA. Note that when using `elsi_ev_real_sparse` and `elsi_ev_complex_sparse`, the eigenvectors are returned in a dense format (BLACS_DENSE), as they are not necessarily sparse.

3.4 Computing Density Matrices

The following subroutines return the density matrix computed from the provided H and S matrices, as well as the energy corresponding to the occupied eigenstates. When the selected solver is ELPA, ELSI will internally construct the density matrix using the eigenvalues and eigenvectors returned by ELPA.

`elsi_dm_real(handle, ham, ovlp, dm, energy)`

Argument	Data Type	in/out	Explanation
<code>handle</code>	<code>type(elsi_handle)</code>	inout	Handle to the current ELSI instance.
<code>ham</code>	real double, rank-2 array	inout	Real Hamiltonian matrix in 2D block-cyclic dense format.
<code>ovlp</code>	real double, rank-2 array	inout	Real overlap matrix (or Cholesky factor) in 2D block-cyclic dense format. See comment 1.
<code>dm</code>	real double, rank-2 array	out	Real density matrix in 2D block-cyclic dense format.
<code>energy</code>	real double	out	Energy corresponding to the occupied eigenstates.

[elsi_dm_complex](#)(handle, ham, ovlp, dm, energy)

Argument	Data Type	in/out	Explanation
handle	type(elsi_handle)	inout	Handle to the current ELSI instance.
ham	complex double, rank-2 array	inout	Complex Hamiltonian matrix in 2D block-cyclic dense format.
ovlp	complex double, rank-2 array	inout	Complex overlap matrix (or its Cholesky factor) in 2D block-cyclic dense format. See comment 1.
dm	complex double, rank-2 array	out	Complex density matrix in 2D block-cyclic dense format.
energy	real double	out	Energy corresponding to the occupied eigenstates.

[elsi_dm_real_sparse](#)(handle, ham, ovlp, dm, energy)

Argument	Data Type	in/out	Explanation
handle	type(elsi_handle)	inout	Handle to the current ELSI instance.
ham	real double, rank-1 array	inout	Non-zero values of the real Hamiltonian matrix in 1D block or block-cyclic CSC format.
ovlp	real double, rank-1 array	inout	Non-zero values of the real overlap matrix in 1D block or block-cyclic CSC format.
dm	real double, rank-1 array	out	Non-zero values of the real density matrix in 1D block or block-cyclic CSC format.
energy	real double	out	Energy corresponding to the occupied eigenstates.

[elsi_dm_complex_sparse](#)(handle, ham, ovlp, dm, energy)

Argument	Data Type	in/out	Explanation
handle	type(elsi_handle)	inout	Handle to the current ELSI instance.
ham	complex double, rank-1 array	inout	Non-zero values of the complex Hamiltonian matrix in 1D block or block-cyclic CSC format.
ovlp	complex double, rank-1 array	inout	Non-zero values of the complex overlap matrix in 1D block or block-cyclic CSC format.
dm	complex double, rank-1 array	out	Non-zero values of the complex density matrix in 1D block or block-cyclic CSC format.
energy	real double	out	Energy corresponding to the occupied eigenstates.

Comments

1) If the chosen solver is ELPA or libOMM, the Hamiltonian matrix will be destroyed during the computation. ELPA will overwrite the overlap matrix in its initial execution with the Cholesky factor, which will be reused by subsequent subroutine calls to [elsi_dm_real](#).

3.5 Customizing ELSI

In ELSI, reasonable default values have been provided for a number of parameters used in the ELSI interface the the supported solvers. However, no set of default parameters can adequately cover all use cases. All parameters can be overridden through ELSI setter subroutines, as described in the following subsections.

3.5.1 Customizing the ELSI Interface

In all the subroutines listed below, the first argument (input and output) is an `elsi_handle`. The second argument (input) of each subroutine is the name of parameter to set.

Note that logical variables are not used in all ELSI API. Integers are used to represent logical, with 0 being false and any positive integer being true.

`elsi_set_output(handle, out_level)`
`elsi_set_write_unit(handle, write_unit)`
`elsi_set_unit_ovlp(handle, unit_ovlp)`
`elsi_set_zero_def(handle, zero_def)`
`elsi_set_sing_check(handle, sing_check)`
`elsi_set_sing_tol(handle, sing_tol)`
`elsi_set_sing_stop(handle, sing_stop)`
`elsi_set_uplo(handle, uplo)`
`elsi_set_mu_broaden_scheme(handle, mu_broaden_scheme)`
`elsi_set_mu_mp_order(handle, mu_mp_order)`
`elsi_set_mu_broaden_width(handle, mu_broaden_width)`
`elsi_set_mu_tol(handle, mu_tol)`

Argument	Data Type	Default	Explanation
<code>out_level</code>	integer	0	Output level of the ELSI interface. 0: no output. 1: standard ELSI output. 2: 1 + info from the solvers. 3: 2 + additional debug info.
<code>write_unit</code>	integer	6	The unit used in ELSI to write out information.
<code>unit_ovlp</code>	integer	0	If not 0, the overlap matrix will be treated as an identity (unit) matrix in ELSI and the solvers. See comment 1.
<code>zero_def</code>	real double	10^{-15}	When converting a matrix from dense to sparse format, values below this threshold will be discarded.
<code>sing_check</code>	integer	0	If not 0, the singularity check of the overlap matrix will be performed. See comment 2.
<code>sing_tol</code>	real double	10^{-5}	Eigenfunctions of the overlap matrix with eigenvalues smaller than this threshold will be removed to avoid ill-conditioning. See comment 1.
<code>sing_stop</code>	integer	0	If not 0, the code always stops if the overlap matrix is detected to be singular. See comment 1.
<code>uplo</code>	integer	0	The input matrix format. 0: Input matrix is full, not triangular. This is currently the only supported option.
<code>mu_broaden_scheme</code>	integer	0	The broadening scheme employed to compute the occupation numbers and the Fermi level. 0: Gaussian. 1: Fermi-Dirac. 2: Methfessel-Paxton. 4: Cold.
<code>mu_mp_order</code>	integer	0	The order of the Methfessel-Paxton broadening scheme. No effect if Methfessel-Paxton is not the chosen broadening scheme.
<code>mu_broaden_width</code>	real double	0.01	The broadening width employed to compute the occupation numbers and the Fermi level. See comment 3.
<code>mu_tol</code>	real double	10^{-13}	The convergence tolerance (in terms of the absolute error in electron count) of the bisection algorithm employed to compute the occupation numbers and the Fermi level.

Comments

1) If the input overlap matrix is an identity matrix, all settings related to the singularity (ill-conditioning) check take no effect.

2) If the singularity check is not disabled, in the first iteration of each SCF cycle, possible singularity of the overlap matrix is checked by computing all its eigenvalues. If there is any eigenvalue smaller than `sing_tol`, the matrix is considered to be singular.

3) In all supported broadening schemes, there is a term $(\epsilon - E_F)/W$ in the distribution function, where ϵ is the energy of an eigenstate, and E_F is the Fermi level. The `broadening_width` parameter should be set to W , in the unit of ϵ and E_F .

3.5.2 Customizing the ELPA Solver

`elsi_set_elpa_solver(handle, elpa_solver)`

Argument	Data Type	Default	Explanation
<code>elpa_solver</code>	integer	2	1: ELPA 1-stage solver. 2: ELPA 2-stage solver. 2 is usually faster and more scalable.

3.5.3 Customizing the libOMM Solver

`elsi_set_omm_flavor(handle, omm_flavor)`

`elsi_set_omm_n_elpa(handle, omm_n_elpa)`

`elsi_set_omm_tol(handle, omm_tol)`

Argument	Data Type	Default	Explanation
<code>omm_flavor</code>	integer	2	Method to perform OMM minimization. See comment 1.
<code>omm_n_elpa</code>	integer	5	Number of ELPA steps before libOMM. See comment 2.
<code>omm_tol</code>	real double	10^{-10}	Convergence tolerance of orbital minimization. See comment 3.

Comments

1) `omm_flavor`: Allowed choices are 0 for basic minimization of a generalized eigenproblem and 2 for a Cholesky factorization of the overlap matrix transforming the generalized eigenproblem to the standard form. Usually 2 leads to a faster convergence of the OMM energy functional minimization, at the price of transforming the eigenproblem. When using sufficiently many steps of ELPA to stabilize the SCF cycle, 1 is a better choice to finish the remaining SCF cycle with possibly higher speed the ELPA or OMM flavor 2. See also comment 2 below.

2) `omm_n_elpa`: It has been demonstrated that OMM is optimal at later stages of an SCF cycle where the electronic structure is closer to its expected local minimum, requiring only one CG iteration to converge the minimization of the OMM energy functional. Accordingly, it is recommended to use ELPA initially, then switching to libOMM after `omm_n_elpa` SCF steps.

3) `omm_tol`: A large minimization tolerance of course leads to a faster convergence, however unavoidably with a lower accuracy. `omm_tol` should be tested and chosen to balance the desired accuracy and computation time of the calling code.

3.5.4 Customizing the PEXSI Solver

`elsi_set_pexsi_n_pole(handle, pexsi_n_pole)`

`elsi_set_pexsi_n_mu(handle, pexsi_n_mu)`

`elsi_set_pexsi_np_per_pole(handle, pexsi_np_per_pole)`

`elsi_set_pexsi_np_symbo(handle, pexsi_np_symbo)`

`elsi_set_pexsi_ordering(handle, pexsi_ordering_method)`

`elsi_set_pexsi_temp(handle, pexsi_temp)`

`elsi_set_pexsi_gap(handle, pexsi_gap)`

`elsi_set_pexsi_delta_e(handle, pexsi_delta_e)`

`elsi_set_pexsi_mu_min(handle, pexsi_mu_min)`

`elsi_set_pexsi_mu_max(handle, pexsi_mu_max)`

`elsi_set_pexsi_inertia_tol(handle, pexsi_inertia_tol)`

Argument	Data Type	Default	Explanation
<code>pexsi_n_pole</code>	integer	20	Number of poles used by PEXSI. See comment 1.
<code>pexsi_n_mu</code>	integer	2	Number of mu points used by PEXSI. See comment 1.
<code>pexsi_np_per_pole</code>	integer	-	Number of MPI tasks assigned to each mu point. See comment 2.
<code>pexsi_np_symbo</code>	integer	1	Number of MPI tasks for symbolic factorization. See comment 3.
<code>pexsi_ordering_method</code>	integer	0	Matrix reordering method. See comment 3.
<code>pexsi_temp</code>	real double	0.002	Temperature. See comment 4.
<code>pexsi_gap</code>	real double	0	Spectral gap. See comment 5.
<code>pexsi_delta_e</code>	real double	10	Spectral radius. See comment 6.
<code>pexsi_mu_min</code>	real double	-10	Minimum value of mu. See comment 7.
<code>pexsi_mu_max</code>	real double	10	Maximum value of mu. See comment 7.
<code>pexsi_inertia_tol</code>	real double	0.05	Stopping criterion of inertia counting. See comment 7.

Comments

1) Starting from PEXSI v1.0.0, the pole expansion is carried out by using the method proposed by Moussa. Usually, 20 poles are enough to get an accuracy that is comparable with the result obtained from diagonalization. Also introduced in PEXSI v1.0.0 is a robust and fast (no longer iterative) algorithm to determine the chemical potential, which performs Fermi operator expansion at several chemical potential values (referred to as “points” by PEXSI developers) in an SCF step, then interpolates the results at all points to get the final answer. The `pexsi_n_mu` setter controls the number of chemical potential “points” to be evaluated. In most cases, 2 points followed by a simple linear interpolation often yield reasonable results.

In short, we recommend `pexsi_n_pole` = 20 and `pexsi_n_mu` = 2.

2) `pexsi_np_per_pole`: PEXSI has, by construction, a 3-level parallelization: the 1st level independently handles all the poles in parallel; within each pole, the 2nd level evaluates the Fermi operator at all the chemical potential points in parallel; finally, within each point, parallel selected inversion is performed in parallel as the 3rd level of parallelization. The value of `pexsi_np_per_pole` is the number of MPI tasks assigned to a single chemical potential point, for the parallel selected inversion at that point. Ideally, the total number of MPI tasks should be `pexsi_np_per_pole` \times `pexsi_n_mu` \times `pexsi_n_pole`, that is, all the three levels of the parallelization in PEXSI are fully exploited. In case that this is not feasible, PEXSI can also process the poles in serial, while the points must be handled simultaneously. As a consequence, the user should make sure that the total number of MPI tasks is a multiple of the number of MPI tasks per point times the number of points. The code will stop if this requirement is not fulfilled.

Assuming the ideal case, the default value for `pexsi_np_per_pole` is: $n_total_mpi / (n_pole \times n_mu)$.

If the dense density matrix solver interfaces are called, the input and output matrices should be 2D-block-cyclic-distributed among all available MPI tasks. In contrast, note that if the sparse density matrix solver interfaces are called, the input and output matrices should be 1D-block-distributed on the first `pexsi_np_per_pole` MPI tasks. PEXSI will then take over and automatically perform any necessary conversions to fit its need.

3) `pexsi_np_symbo`: The number of MPI tasks used for symbolic factorization cannot be larger than a magic number, depending on the matrix and the machine. The default value is set to 1 for safety. It is worth testing and increasing this number. The default matrix reordering method in PEXSI is set to parallel ordering using PtScotch or ParMETIS (0). The sequential ordering (1) may be necessary in rare cases.

4) `pexsi_temp`: This value corresponds to the $1/k_B T$ term in the Fermi-Dirac distribution function.

5) `pexsi_gap`: The PEXSI method does not require a gap. If an estimate of the gap is unavailable, the default value usually works.

6) `pexsi.delta_e`: This is the spectral width of the eigensystem, i.e., the difference between the largest and smallest eigenvalues. Use the default value if no access to a better estimate.

7) The newly designed chemical potential determination algorithm relies on inertia counting to narrow down the chemical potential searching interval in the first few SCF steps. The `pexsi.inertia_tol` setter controls the stopping criterion of the inertia counting procedure. With a small interval obtained from the inertia counting step, PEXSI then selects a number of points in this interval to perform Fermi operator calculations, based on which a final chemical potential will be determined. The trick of this algorithm is that the chemical potential interval of the current SCF step can be used as a descent guess in the next SCF step. Therefore, the mechanism to choose input values for `pexsi.mu_min` and `pexsi.mu_max` is two-fold. For the first SCF iteration, they should be set to safe values that guarantee the true chemical potential lies in this interval. Then, for the n^{th} SCF step, `pexsi.mu_min` should be set to $(\mu_{min}^{n-1} + \Delta V_{min})$, `pexsi.mu_max` should be set to $(\mu_{max}^{n-1} + \Delta V_{max})$. Here, μ_{min}^{n-1} and μ_{max}^{n-1} are the lower bound and the upper bound of the chemical potential that are determined by PEXSI in the $(n-1)^{th}$ SCF step. They can be retrieved by calling `elsi.get_pexsi_mu_min` and `elsi.get_pexsi_mu_max`, respectively (see Section 3.6.2). Suppose the effective potential (Hartree potential, exchange-correlation potential, and external potential) is stored in an array V , whose dimension is the number of grid points. From one SCF iteration to the next, ΔV denotes the potential change, and ΔV_{min} and ΔV_{max} are the minimum and maximum values in the array ΔV , respectively. The whole process is summarized in the following pseudo-code.

```
mu_min = -10.0
mu_max = 10.0
 $\Delta V_{min}$  = 0.0
 $\Delta V_{max}$  = 0.0

while SCF not converged do
    Update Hamiltonian

    elsi_set_pexsi_mu_min(elsi_h, mu_min +  $\Delta V_{min}$ )
    elsi_set_pexsi_mu_max(elsi_h, mu_max +  $\Delta V_{max}$ )

    elsi_dm_{real|complex}(elsi_h, ham, ovlp, dm, energy)

    elsi_get_pexsi_mu_min(elsi_h, mu_min)
    elsi_get_pexsi_mu_max(elsi_h, mu_max)

    Update electron density
    Update potential

     $\Delta V_{min}$  = minval( $V_{new} - V_{old}$ )
     $\Delta V_{max}$  = maxval( $V_{new} - V_{old}$ )

    Check SCF convergence
end
```

3.6 Getting Additional Results from ELSI

In Section 3.3 and Section 3.4, the interfaces to compute and return the eigensolutions and the density matrices have been introduced. Internally, ELSI and the solvers perform additional calculations whose results may not necessarily be useful for a normal SCF calculation, but can be valuable in some cases. One example is the energy-weighted density matrix that is employed to evaluate the Pulay forces during a geometry optimization calculation. The ELSI getter subroutines are used to retrieve such additional results from ELSI, as will be described in the following subsections.

3.6.1 Getting Results from the ELSI Interface

In all the subroutines listed below, the first argument (input and output) is an `elsi_handle`. The second argument (output) of each subroutine is the name of parameter to get.

```
elsi_get_n_sing(handle, n_sing)
elsi_get_mu(handle, mu)
elsi_get_entropy(handle, ts)
elsi_get_edm_real(handle, edm_real)
elsi_get_edm_complex(handle, edm_complex)
elsi_get_edm_real_sparse(handle, edm_real_sparse)
elsi_get_edm_complex_sparse(handle, edm_complex_sparse)
```

Argument	Data Type	Explanation
n_sing	integer	Number of eigenvalues of the overlap matrix that are smaller than the singularity tolerance. See Section 3.5.1.
mu	real double	Chemical potential. See comment 1.
ts	real double	Entropy. See comment 1.
edm_real	real double, rank-2 array	Real energy-weighted density matrix in 2D block-cyclic dense format. See comment 2.
edm_complex	complex double, rank-2 array	Complex energy-weighted density matrix in 2D block-cyclic dense format. See comment 2.
edm_real_sparse	real double, rank-1 array	Non-zero values of the real density matrix in 1D block CSC format. See comment 2.
edm_complex_sparse	complex double, rank-1 array	Non-zero values of the complex density matrix in 1D block CSC format. See comment 2.

Comments

1) In ELSI, the chemical potential will be available if one of the density matrix solver interfaces has been called, with either ELPA or PEXSI being the chosen solver. Then, the chemical potential can be retrieved by calling [elsi_get_mu](#). Similarly, the entropy will be available if one of the density matrix solver interfaces has been called with ELPA being the chosen solver. The user should avoid calling the subroutine when the chemical potential is not ready. The code will print a warning in such a case.

2) In general, the energy-weighted density matrix is only needed in a late stage of an SCF cycle to evaluate forces. It is, therefore, not calculated when any of the density matrix solver interface is called. When the energy-weighted density matrix is actually needed, it can be requested by calling the [elsi_get_edm](#) subroutines. Note that these subroutines all have the requirement that the corresponding [elsi_dm](#) subroutine must have been invoked. For instance, [elsi_get_edm_real_sparse](#) only makes sense if [elsi_dm_real_sparse](#) has been successfully executed.

3.6.2 Getting Results from the PEXSI Solver

```
elsi_get_pexsi_mu_min(handle, pexsi_mu_min)
elsi_get_pexsi_mu_max(handle, pexsi_mu_max)
```

Argument	Data Type	Explanation
pexsi_mu_min	real double	Minimum value of mu. See comment 1.
pexsi_mu_max	real double	Maximum value of mu. See comment 1.

Comments

1) Please refer to the 7th comment in Section 3.5.4 for the chemical potential determination algorithm implemented in PEXSI and ELSI.

3.7 Demonstration Pseudo-Code

The typical workflow of ELSI within a KS-DFT code is demonstrated by the following pseudo-code. The ELSI interface routines and the operations to be performed by a KS-DFT code are marked with blue and red, respectively.

3.7.1 2D Block-Cyclic Distributed Dense Matrix + ELSI Eigensolver Interface

SCF initialize

```
elsi_init(elsi_h, ELPA, MULTIPROC, BLACS_DENSE, n_basis, n_electron, n_state)
elsi_set_mpi(elsi_h, mpi_comm)
elsi_set_blacs(elsi_h, blacs_ctxt, block_size)

while SCF not converged do
    Update Hamiltonian

    elsi_ev_{real|complex}(elsi_h, ham, ovlp, eval, evec)

    Update electron density
    Check SCF convergence
end

elsi_finalize(elsi_h)
```

3.7.2 1D Block Distributed CSC Sparse Matrix + ELSI Eigensolver Interface

SCF initialization

```
elsi_init(elsi_h, ELPA, MULTIPROC, PEXSL_CSC, n_basis, n_electron, n_state)
elsi_set_mpi(elsi_h, mpi_comm)
elsi_set_blacs(elsi_h, blacs_ctxt, block_size)
elsi_set_csc(elsi_h, global_nnz, local_nnz, local_col, row_index, col_pointer)

while SCF not converged do
    Update Hamiltonian

    elsi_ev_{real|complex}_sparse(elsi_h, ham, ovlp, eval, evec)

    Update electron density
    Check SCF convergence
end

elsi_finalize(elsi_h)
```

Comments

- 1) The SINGLE_PROC parallel mode is not compatible with the PEXSL_CSC matrix format.
- 2) Although the input Hamiltonian and overlap matrices are in a sparse format, the calculated eigenvectors are not in the same sparsity pattern. In fact, they are usually not sparse. Therefore, the eigenvectors are returned in the

BLACS_DENSE format, which is required to be properly set up before calling `elsi_ev_{real|complex}`.

3.7.3 1D Block-Cyclic Distributed CSC Sparse Matrix + ELSI Eigensolver Interface

SCF initialization

```
elsi_init(elsi_h, ELPA, MULTI_PROC, SIESTA_CSC, n_basis, n_electron, n_state)
elsi_set_mpi(elsi_h, mpi_comm)
elsi_set_blacs(elsi_h, blacs_ctxt, block_size)
elsi_set_csc(elsi_h, global_nnz, local_nnz, local_col, row_index, col_pointer)
elsi_set_csc_blk(elsi_h, block_size)

while SCF not converged do
    Update Hamiltonian

    elsi_ev_{real|complex}_sparse(elsi_h, ham, ovlp, eval, evec)

    Update electron density
    Check SCF convergence
end

elsi_finalize(elsi_h)
```

Comments

- 1) The SINGLE_PROC parallel mode is not compatible with the SIESTA_CSC matrix format.
- 2) Although the input Hamiltonian and overlap matrices are in a sparse format, the calculated eigenvectors are not in the same sparsity pattern. In fact, they are usually not sparse. Therefore, the eigenvectors are returned in the BLACS_DENSE format, which is required to be properly set up before calling `elsi_ev_{real|complex}`.

3.7.4 2D Block-Cyclic Distributed Dense Matrix + ELSI Density Matrix Interface

SCF initialization

```
elsi_init(elsi_h, LIBOMM, MULTIPROC, BLACS_DENSE, n_basis, n_electron, n_state)
elsi_set_mpi(elsi_h, mpi_comm)
elsi_set_blacs(elsi_h, blacs_ctxt, block_size)

while SCF not converged do
    Update Hamiltonian

    elsi_dm_{real|complex}(elsi_h, ham, ovlp, dm, energy)
    elsi_get_edm_{real|complex}(elsi_h, edm)

    Update electron density
    Check SCF convergence
end

elsi_finalize(elsi_h)
```

Comments

- 1) The SINGLE_PROC parallel mode is not compatible with the ELSI density matrix interface.
- 2) The returned energy corresponds to the eigenenergies of the occupied Kohn-Sham orbitals.

3.7.5 1D Block Distributed CSC Sparse Matrix + ELSI Density Matrix Interface

SCF initialization

```
elsi_init(elsi_h, PEXSI, parallel_mode, PEXSI_CSC, n_basis, n_electron, n_state)
elsi_set_mpi(elsi_h, mpi_comm)
elsi_set_csc(elsi_h, global_nnz, local_nnz, local_col, row_index, col_pointer)

while SCF not converged do
    Update Hamiltonian

    elsi_dm_{real|complex}_sparse(elsi_h, ham, ovlp, dm, energy)
    elsi_get_edm_{real|complex}_sparse(elsi_h, edm)

    Update electron density
    Check SCF convergence
end

elsi_finalize(elsi_h)
```

Comments

- 1) The SINGLE_PROC parallel mode is not compatible with the ELSI density matrix interface.

- 2) Refer to the 7th comment in Section 3.5.4 for the chemical potential determination algorithm implemented in PEXSI.
- 3) The returned energy corresponds to the eigenenergies of the occupied Kohn-Sham orbitals.

3.7.6 1D Block-Cyclic Distributed CSC Sparse Matrix + ELSI Density Matrix Interface

SCF initialization

```
elsi_init(elsi_h, PEXSI, parallel_mode, SIESTA_CSC, n_basis, n_electron, n_state)
elsi_set_mpi(elsi_h, mpi_comm)
elsi_set_csc(elsi_h, global_nnz, local_nnz, local_col, row_index, col_pointer)
elsi_set_csc_blk(elsi_h, block_size)
```

```
while SCF not converged do
```

```
    Update Hamiltonian
```

```
    elsi_dm_{real|complex}_sparse(elsi_h, ham, ovlp, dm, energy)
    elsi_get_edm_{real|complex}_sparse(elsi_h, edm)
```

```
    Update electron density
```

```
    Check SCF convergence
```

```
end
```

```
elsi_finalize(elsi_h)
```

Comments

- 1) The SINGLE_PROC parallel mode is not compatible with the ELSI density matrix interface.
- 2) Refer to the 7th comment in Section 3.5.4 for the chemical potential determination algorithm implemented in PEXSI.
- 3) The returned energy corresponds to the eigenenergies of the occupied Kohn-Sham orbitals.

3.7.7 Multiple k-points Calculations

SCF initialization

```
elsi_init(elsi_h, ELPA, parallel_mode, BLACS_DENSE, n_basis, n_electron, n_state)
elsi_set_mpi(elsi_h, mpi_comm)
elsi_set_blacs(elsi_h, blacs_ctxt, block_size)

elsi_set_kpoint(elsi_h, n_kpt, i_kpt, weight)
elsi_set_mpi_global(elsi_h, mpi_comm_global)

while SCF not converged do
    Update Hamiltonian

    elsi_dm_{real|complex}(elsi_h, ham, ovlp, dm, energy)
    elsi_get_edm_{real|complex}(elsi_h, edm)

    Update electron density
    Check SCF convergence
end

elsi_finalize(elsi_h)
```

Comments

- 1) When there are multiple k-points, other than setting up the k-points and a global MPI communicator, there is no change in the way ELSI solver interfaces are called.
- 2) Calculations with two spin channels can be set up similarly.
- 3) One density matrix will be returned for each k-point. The KS-DFT code will take over to assemble the real-space density matrix. The returned energy, however, is already summed over all k-points with respect to the weight of each k-point.

3.8 C/C++ Interface

ELSI is written in Fortran, but many KS-DFT codes use C/C++. ELSI provides a C interface around the core Fortran ELSI code, which can be called from a C or C++ program. Each C wrapper function corresponds to a Fortran subroutine, where we have prefixed the original Fortran subroutine name with `c_` for clarity and consistency. Argument lists are identical to the associated native Fortran subroutine. For the complete definition of the C interface, the user is encouraged to look at the `elsi.h` header file directly.

Bibliography

- [1] W. Kohn and L.J. Sham, Self-consistent equations including exchange and correlation effects, *Physical Review*, 140, 1133-1138 (1965).
- [2] F. Corsetti, The orbital minimization method for electronic structure calculations with finite-range atomic basis sets, *Computer Physics Communications*, 185, 873-883 (2014).
- [3] T. Auckenthaler et al., Parallel solution of partial symmetric eigenvalue problems from electronic structure calculations, *Parallel Computing*, 37, 783-794 (2011).
- [4] A. Marek et al., The ELPA library: Scalable parallel eigenvalue solutions for electronic structure theory and computational science, *Journal of Physics: Condensed Matter*, 26, 213201 (2014).
- [5] L. Lin et al., Fast algorithm for extracting the diagonal of the inverse matrix with application to the electronic structure analysis of metallic systems, *Communications in Mathematical Sciences*, 7, 755-777 (2009).
- [6] L. Lin et al., Accelerating atomic orbital-based electronic structure calculation via pole expansion and selected inversion, *Journal of Physics: Condensed Matter*, 25, 295501 (2013).
- [7] B. Aradi et al., DFTB+, a sparse matrix-based implementation of the DFTB method, *Journal of Physical Chemistry A*, 111, 5678 (2007).
- [8] W. Hu et al., DGDFT: A massively parallel method for large scale density functional theory calculations, *The Journal of Chemical Physics*, 143, 124110 (2015).
- [9] V. Blum et al., Ab initio molecular simulations with numeric atom-centered orbitals, *Computer Physics Communications*, 180, 2175-2196 (2009).
- [10] M. Valiev et al., NWChem: A comprehensive and scalable open-source solution for large scale molecular simulations, *Computer Physics Communications*, 181, 1477-1489 (2010).
- [11] J.M. Soler et al., The SIESTA method for ab initio order-N materials simulation, *Journal of Physics: Condensed Matter*, 14, 2745-2779 (2002).
- [12] V. Yu et al., ELSI: A unified software interface for Kohn-Sham electronic structure solvers, *Computer Physics Communications*, <http://dx.doi.org/10.1016/j.cpc.2017.09.007>.

License and Copyright

ELSI Interface software is licensed under the 3-clause BSD license:

Copyright (c) 2015-2018, the ELSI team. All rights reserved.

Redistribution and use in source and binary forms, with or without modification, are permitted provided that the following conditions are met:

- 1) Redistributions of source code must retain the above copyright notice, this list of conditions and the following disclaimer.
- 2) Redistributions in binary form must reproduce the above copyright notice, this list of conditions and the following disclaimer in the documentation and/or other materials provided with the distribution.
- 3) Neither the name of the "Electronic Structure Infrastructure (ELSI)" project nor the names of its contributors may be used to endorse or promote products derived from this software without specific prior written permission.

THIS SOFTWARE IS PROVIDED BY THE COPYRIGHT HOLDERS AND CONTRIBUTORS "AS IS" AND ANY EXPRESS OR IMPLIED WARRANTIES, INCLUDING, BUT NOT LIMITED TO, THE IMPLIED WARRANTIES OF MERCHANTABILITY AND FITNESS FOR A PARTICULAR PURPOSE ARE DISCLAIMED. IN NO EVENT SHALL COPYRIGHT HOLDER BE LIABLE FOR ANY DIRECT, INDIRECT, INCIDENTAL, SPECIAL, EXEMPLARY, OR CONSEQUENTIAL DAMAGES (INCLUDING, BUT NOT LIMITED TO, PROCUREMENT OF SUBSTITUTE GOODS OR SERVICES; LOSS OF USE, DATA, OR PROFITS; OR BUSINESS INTERRUPTION) HOWEVER CAUSED AND ON ANY THEORY OF LIABILITY, WHETHER IN CONTRACT, STRICT LIABILITY, OR TORT (INCLUDING NEGLIGENCE OR OTHERWISE) ARISING IN ANY WAY OUT OF THE USE OF THIS SOFTWARE, EVEN IF ADVISED OF THE POSSIBILITY OF SUCH DAMAGE.

The source code of ELPA (LGPL), libOMM (2-clause BSD), PEXSI (3-clause BSD), and SuperLU_DIST (3-clause BSD) are redistributed within ELSI. Individual license of each library can be found in the corresponding subfolder.